

DeepLearning1_Supervised_mbochenek

May 16, 2022

1 1000-719bMSB Modeling of Complex Biological Systems

2 Deep Neural Network: Supervised Learning

2.0.1 Basic python and pandas

<https://www.kaggle.com/lavanyashukla01/pandas-numpy-python-cheatsheet>

<https://www.utc.fr/~jlaforet/Suppl/python-cheatsheets.pdf>

List comprehensions are a concise way to create new lists from existing ones.

```
[1]: list1 = list(range(0,10))  
      print(list1)
```

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
```

```
[2]: list1[0] # A vector in R starts with an index of 1. In Python, 0.
```

```
[2]: 0
```

```
[3]: list1[2:5]
```

```
[3]: [2, 3, 4]
```

```
[4]: list2 = []  
      for i in list1:  
          list2.append(i+1)  
  
      print(list2)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

```
[5]: list3 = [i+1 for i in list1]  
  
      print(list3)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

3 Classification of MNIST using densely connected layers

We are going to use the Keras library to implement a neural network that can classify handwritten digits - in just a few lines of code.

First we load and inspect the data. The dataset is split into training and test data.

```
[6]: import numpy as np
import tensorflow.keras as keras
import matplotlib.pyplot as plt
```

```
[7]: import tensorflow as tf
print(tf.__version__)
tf.compat.v1.disable_eager_execution()
```

2.8.0

```
[8]: (train_images, train_labels), (test_images, test_labels) = keras.datasets.mnist.
    ↪load_data()
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11493376/11490434 [=====] - 0s 0us/step

11501568/11490434 [=====] - 0s 0us/step

```
[9]: train_images.shape
```

```
[9]: (60000, 28, 28)
```

```
[10]: train_labels.shape
```

```
[10]: (60000,)
```

```
[11]: test_images.shape
```

```
[11]: (10000, 28, 28)
```


```
[12]: test_labels.shape
```

```
[12]: (10000,)
```

Let's plot one of the digits and the corresponding label.

```
[16]: print('Label of element 0:', train_labels[1])
plt.matshow(train_images[1], cmap='gray')
plt.show()
```

Label of element 0: 0



In this step we define the neural network. ReLu is an activation function defined as $f(x) = \max(0, x)$. Softmax activation function is normalized such that the sum of all outputs is equal 1.

```
[17]: from tensorflow.keras import layers
      from tensorflow.keras import models
      model = models.Sequential()
      model.add(layers.Dense(512, activation='relu', input_shape=(28 * 28,)))
      #units (here, 512) - Positive integer, dimensionality of the output space.
      model.add(layers.Dense(10, activation='softmax'))
```

With compile we tell the network which optimizer and loss function to use. Optimizer specifies the particular implementation of the gradient-descent, e.g. how it adapts the learning rate. 'Metrics' specifies the output during the training.

```
[18]: model.compile(optimizer='rmsprop',
                  loss='mean_squared_error',
                  metrics=['accuracy'])
```

```
[19]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 512)	401920

dense_1 (Dense) (None, 10) 5130

```
=====
Total params: 407,050
Trainable params: 407,050
Non-trainable params: 0
-----
```

We are using a densely connected network, so we have to flatten the images.
Input values should be in the range (0,1) for fast convergence.

```
[21]: #reshaping training dataset
train_images_flat = train_images.reshape((60000, 28 * 28))
#normalization - finding max value in the train dataset
maxValue = np.amax(train_images_flat)
print(maxValue)
#now normalizing
train_images_flat = train_images_flat.astype('float32') / 255

#reshaping testing dataset
test_images_flat = test_images.reshape((10000, 28 * 28))
#normalization - dividing by max value
test_images_flat = test_images_flat.astype('float32') / 255
```

255

Convert the labels to a 'one-hot' coding.

```
[22]: from tensorflow.keras.utils import to_categorical
print(f"Train labels before 'one hot coding': {train_labels[:10]}")
train_labels = to_categorical(train_labels)
print(f"Train labels after 'one hot coding': {train_labels[:10]}")
test_labels = to_categorical(test_labels)
```

```
Train labels before 'one hot coding': [5 0 4 1 9 2 1 3 1 4]
Train labels after 'one hot coding': [[0. 0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]
 [0. 0. 1. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]
```


```
[23]: train_images.reshape((60000,28*28)).shape
```

```
[23]: (60000, 784)
```

```
[24]: train_labels[0]
```

```
[24]: array([0., 0., 0., 0., 0., 1., 0., 0., 0., 0.], dtype=float32)
```

```
[25]: plt.matshow(train_images[0], cmap='gray')  
plt.show()
```



```
[26]: model.fit(train_images_flat, train_labels, epochs=5, batch_size=128)
```

Train on 60000 samples

Epoch 1/5

60000/60000 [=====] - 4s 61us/sample - loss: 0.0121 -
accuracy: 0.9226

Epoch 2/5

60000/60000 [=====] - 1s 19us/sample - loss: 0.0053 -
accuracy: 0.9671

Epoch 3/5

60000/60000 [=====] - 1s 19us/sample - loss: 0.0037 -
accuracy: 0.9771

Epoch 4/5

60000/60000 [=====] - 1s 19us/sample - loss: 0.0028 -
accuracy: 0.9826

Epoch 5/5

```
60000/60000 [=====] - 1s 19us/sample - loss: 0.0022 -  
accuracy: 0.9861
```

```
[26]: <keras.callbacks.History at 0x7f8633576150>
```

Let's check the performance on the test set. If the accuracy is less than the training accuracy, then we might be overfitting!

```
[27]: test_loss, test_acc = model.evaluate(test_images_flat, test_labels)  
print('test_acc:', test_acc)
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2057:  
UserWarning: `Model.state_updates` will be removed in a future version. This  
property should not be used in TensorFlow 2.0, as `updates` are applied  
automatically.
```

```
    updates = self.state_updates
```

```
test_acc: 0.9805
```

We can also find the predictions for a selection of input images.


```
[28]: predictions = model.predict(train_images_flat[:10])
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2079:  
UserWarning: `Model.state_updates` will be removed in a future version. This  
property should not be used in TensorFlow 2.0, as `updates` are applied  
automatically.
```

```
    updates=self.state_updates,
```

```
[31]: img_num = 9  
print(predictions[img_num])  
print(train_labels[img_num])  
plt.matshow(train_images[img_num], cmap='gray')  
plt.show()
```

```
[1.8075972e-06 3.0627989e-09 2.8879535e-06 9.8233635e-08 9.9992228e-01  
 7.7069035e-06 6.0323819e-05 4.0500558e-06 3.0447688e-07 5.6681574e-07]  
[0. 0. 0. 0. 1. 0. 0. 0. 0. 0.]
```



```
[32]: train_labels
```

```
[32]: array([[0., 0., 0., ..., 0., 0., 0.],
            [1., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            ...,
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 0., 0.],
            [0., 0., 0., ..., 0., 1., 0.]], dtype=float32)
```

4 Classification of MNIST using convolutional layers

We have build a classifier for handwritten images only using densely connected layers. Let's see if we can do better using convolutional layers!

First define the convolutional layers.

```
[33]: model2 = models.Sequential()
model2.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
model2.add(layers.MaxPooling2D((2, 2)))
model2.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
[34]: model2.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928

=====
Total params: 55,744
Trainable params: 55,744
Non-trainable params: 0
=====

Now add a classifier on top of the convnet.

```
[35]: model2.add(layers.Flatten()) #flattening layer
model2.add(layers.Dense(64, activation='relu')) #dense layer
model2.add(layers.Dense(10, activation='softmax')) #dense layer
```

```
[36]: model2.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928

flatten (Flatten)	(None, 576)	0
dense_2 (Dense)	(None, 64)	36928
dense_3 (Dense)	(None, 10)	650

```

=====
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
-----

```

```
[37]: train_images_conv = train_images.reshape((60000, 28, 28, 1))
      train_images_conv = train_images_conv.astype('float32') / 255
      test_images_conv = test_images.reshape((10000, 28, 28, 1))
      test_images_conv = test_images_conv.astype('float32') / 255
```

```
[38]: model2.compile(optimizer='rmsprop',
                    loss='categorical_crossentropy',
                    metrics=['accuracy'])
```

```
[39]: model2.fit(train_images_conv, train_labels, epochs=15, batch_size=64)
```

```

Train on 60000 samples
Epoch 1/15
60000/60000 [=====] - 13s 220us/sample - loss: 0.1779 -
accuracy: 0.9438
Epoch 2/15
60000/60000 [=====] - 4s 74us/sample - loss: 0.0479 -
accuracy: 0.9851
Epoch 3/15
60000/60000 [=====] - 4s 64us/sample - loss: 0.0330 -
accuracy: 0.9899
Epoch 4/15
60000/60000 [=====] - 3s 57us/sample - loss: 0.0253 -
accuracy: 0.9921
Epoch 5/15
60000/60000 [=====] - 4s 74us/sample - loss: 0.0193 -
accuracy: 0.9941
Epoch 6/15
60000/60000 [=====] - 4s 66us/sample - loss: 0.0158 -
accuracy: 0.9950
Epoch 7/15
60000/60000 [=====] - 4s 68us/sample - loss: 0.0132 -
accuracy: 0.9961
Epoch 8/15
60000/60000 [=====] - 4s 61us/sample - loss: 0.0098 -
accuracy: 0.9969

```

```

Epoch 9/15
60000/60000 [=====] - 4s 60us/sample - loss: 0.0099 -
accuracy: 0.9971
Epoch 10/15
60000/60000 [=====] - 4s 61us/sample - loss: 0.0086 -
accuracy: 0.9978
Epoch 11/15
60000/60000 [=====] - 4s 60us/sample - loss: 0.0069 -
accuracy: 0.9980
Epoch 12/15
60000/60000 [=====] - 4s 65us/sample - loss: 0.0069 -
accuracy: 0.9981
Epoch 13/15
60000/60000 [=====] - 4s 64us/sample - loss: 0.0058 -
accuracy: 0.9984
Epoch 14/15
60000/60000 [=====] - 3s 57us/sample - loss: 0.0049 -
accuracy: 0.9987
Epoch 15/15
60000/60000 [=====] - 4s 59us/sample - loss: 0.0041 -
accuracy: 0.9989

```

[39]: <keras.callbacks.History at 0x7f85b45d6f10>

```

[40]: test_loss, test_acc = model2.evaluate(test_images_conv, test_labels)
      print(f"Accuracy on the test dataset = {test_acc}")

```

```

/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2057:
UserWarning: `Model.state_updates` will be removed in a future version. This
property should not be used in TensorFlow 2.0, as `updates` are applied
automatically.

```

```

    updates = self.state_updates

```

```

Accuracy on the test dataset = 0.9925000071525574

```

4.1 Introducing Fashion MNIST (Homework dataset)

The MNIST dataset is not too demanding, let's try something a little more difficult - Fashion MNIST.

[LINK TO IMAGE](#)

Check out labels on [GitHub](#)

```

[63]: (train_imgs_fash, train_labels_fash), (test_imgs_fash, test_labels_fash) = ↵
      ↪keras.datasets.fashion_mnist.load_data()


```

```
[47]: print(train_imgs_fash.shape)
      print(test_imgs_fash.shape)
```

```
(60000, 28, 28)
(10000, 28, 28)
```

```
[48]: plt.imshow(train_imgs_fash[12], cmap=plt.get_cmap('gray'))
```

```
[48]: <matplotlib.image.AxesImage at 0x7f85b43a5d10>
```



5 HOMEWORK 1

Build a classifier for fashion MNIST.

1. Use exactly the same architectures (both densely connected layers and from convolutional layers) as the above MNIST e.g., replace the dataset. Save the Jupyter Notebook in its original format and output a PDF file after training, testing, and validation. Make sure to write down how do they perform (training accuracy, testing accuracy).

2. Improve the architecture. Experiment with different numbers of layers, size of layers, number of filters, size of filters. You are required to make those adjustment to get the highest accuracy. Watch out for overfitting -- we want the highest testing accuracy! Please provide a PDF file of the result, the best test accuracy and the architecture (different numbers of layers, size of layers, number of filters, size of filters)

5.0.1 Data preprocessing

```
[64]: #reshaping
train_imgs_fash = train_imgs_fash.reshape((60000, 28, 28, 1))
test_imgs_fash = test_imgs_fash.reshape((10000, 28, 28, 1))

#normalizing data
#finding max value in the dataset
maxValue = np.amax(train_imgs_fash)
print(f"max value = {maxValue}")

train_imgs_fash = train_imgs_fash.astype('float32') / maxValue
test_imgs_fash = test_imgs_fash.astype('float32') / maxValue

#encoding labels
print(f"Train labels before 'one hot coding': {train_labels_fash[:10]}")
train_labels_fash = to_categorical(train_labels_fash)
print(f"Train labels after 'one hot coding': {train_labels_fash[:10]}")
test_labels_fash = to_categorical(test_labels_fash)
```

max value = 255

Train labels before 'one hot coding': [9 0 0 3 0 2 7 2 5 5]

Train labels after 'one hot coding': [[0. 0. 0. 0. 0. 0. 0. 0. 0. 1.]

[1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

[1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

[0. 0. 0. 1. 0. 0. 0. 0. 0. 0.]

[1. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

[0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]

[0. 0. 0. 0. 0. 0. 0. 1. 0. 0.]

[0. 0. 1. 0. 0. 0. 0. 0. 0. 0.]

[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

[0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

5.0.2 First model creation

```
[53]: #creating the model
model_fash1 = models.Sequential()
model_fash1.add(layers.Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)))
model_fash1.add(layers.MaxPooling2D((2, 2)))
model_fash1.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_fash1.add(layers.MaxPooling2D((2, 2)))
model_fash1.add(layers.Conv2D(64, (3, 3), activation='relu'))
model_fash1.add(layers.Flatten()) #flattening layer
model_fash1.add(layers.Dense(64, activation='relu')) #dense layer
model_fash1.add(layers.Dense(10, activation='softmax')) #dense layer
```

```
#compiling
model_fash1.compile(optimizer='rmsprop',
loss='categorical_crossentropy',
metrics=['accuracy'])
```

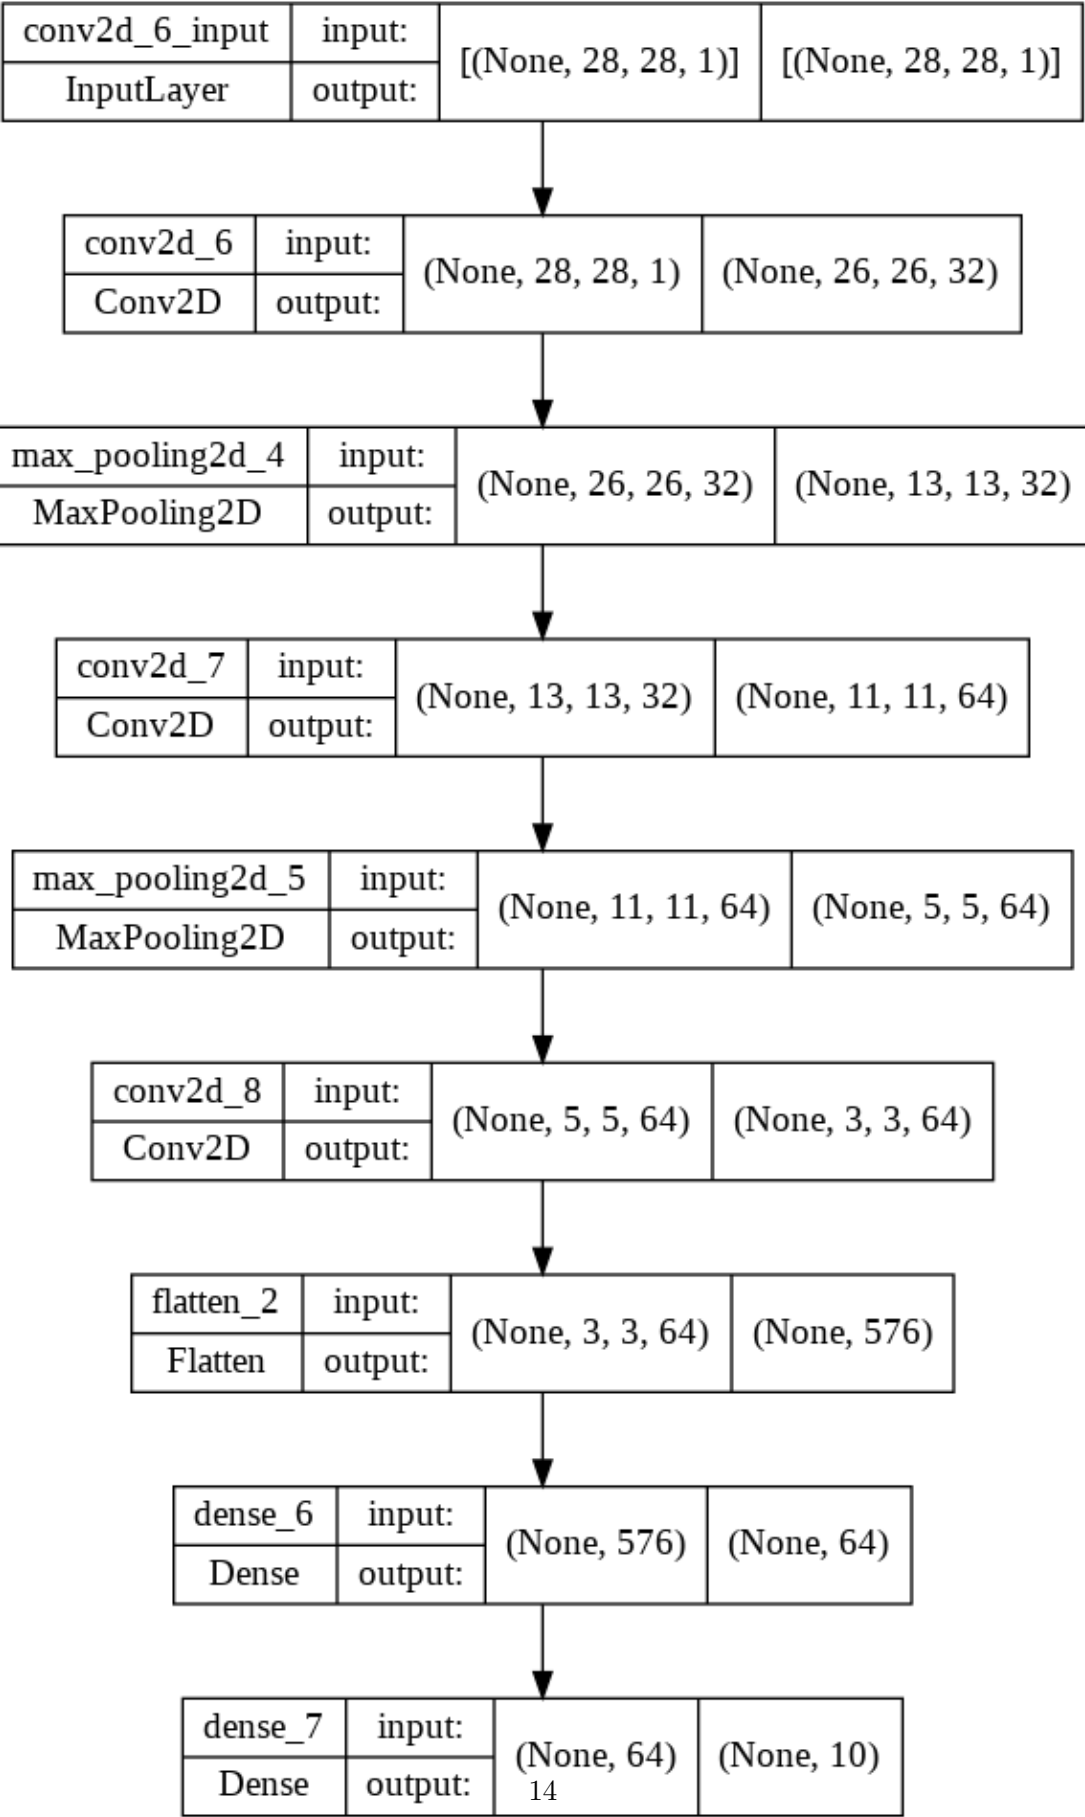
```
[59]: model_fash1.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d_4 (MaxPooling 2D)	(None, 13, 13, 32)	0
conv2d_7 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_5 (MaxPooling 2D)	(None, 5, 5, 64)	0
conv2d_8 (Conv2D)	(None, 3, 3, 64)	36928
flatten_2 (Flatten)	(None, 576)	0
dense_6 (Dense)	(None, 64)	36928
dense_7 (Dense)	(None, 10)	650
Total params: 93,322		
Trainable params: 93,322		
Non-trainable params: 0		

```
[147]: tf.keras.utils.plot_model(model_fash1, 'model_fash1.png', show_shapes=True)
```

```
[147]:
```



```
[54]: model_fash1.fit(train_imgs_fash, train_labels_fash, epochs=15, batch_size=64)
```

Train on 60000 samples

Epoch 1/15

60000/60000 [=====] - 4s 62us/sample - loss: 0.5396 -
accuracy: 0.8001

Epoch 2/15

60000/60000 [=====] - 4s 62us/sample - loss: 0.3333 -
accuracy: 0.8794

Epoch 3/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.2788 -
accuracy: 0.8985

Epoch 4/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.2480 -
accuracy: 0.9093

Epoch 5/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.2249 -
accuracy: 0.9168

Epoch 6/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.2072 -
accuracy: 0.9239

Epoch 7/15

60000/60000 [=====] - 4s 59us/sample - loss: 0.1902 -
accuracy: 0.9300

Epoch 8/15

60000/60000 [=====] - 4s 59us/sample - loss: 0.1760 -
accuracy: 0.9355

Epoch 9/15

60000/60000 [=====] - 4s 58us/sample - loss: 0.1645 -
accuracy: 0.9396

Epoch 10/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.1530 -
accuracy: 0.9427

Epoch 11/15

60000/60000 [=====] - 3s 57us/sample - loss: 0.1429 -
accuracy: 0.9479

Epoch 12/15

60000/60000 [=====] - 3s 57us/sample - loss: 0.1337 -
accuracy: 0.9507

Epoch 13/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.1256 -
accuracy: 0.9536

Epoch 14/15

60000/60000 [=====] - 3s 58us/sample - loss: 0.1179 -
accuracy: 0.9563

```
Epoch 15/15
60000/60000 [=====] - 3s 58us/sample - loss: 0.1118 -
accuracy: 0.9585
```

```
[54]: <keras.callbacks.History at 0x7f85b4212a90>
```

```
[55]: test_loss, test_acc = model_fash1.evaluate(test_imgs_fash, test_labels_fash)
      print(f"Accuracy on the test dataset = {test_acc}")
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2057:
UserWarning: `Model.state_updates` will be removed in a future version. This
property should not be used in TensorFlow 2.0, as `updates` are applied
automatically.
```

```
    updates = self.state_updates
```

```
Accuracy on the test dataset = 0.9110000133514404
```

5.0.3 Improve the architecture

```
[127]: #creating the model
model_fash2 = models.Sequential()
model_fash2.add(layers.Conv2D(48, (3, 3), padding='same', activation='relu',
    ↪input_shape=(28, 28, 1)))
model_fash2.add(layers.MaxPooling2D(pool_size=(2, 2), strides=2,
    padding='valid'))
model_fash2.add(layers.Conv2D(96, (3, 3), padding='same', activation='relu'))
model_fash2.add(layers.MaxPooling2D(pool_size=(2, 2), strides=2,
    padding='valid'))
model_fash2.add(layers.Conv2D(96, (3, 3), padding='same', activation='relu'))

model_fash2.add(layers.Flatten()) #flattening layer
model_fash2.add(layers.Dense(128, activation='relu')) #dense layer
model_fash2.add(layers.Dense(10, activation='softmax')) #dense layer

#compiling
model_fash2.compile(optimizer='rmsprop',
loss='categorical_crossentropy',
metrics=['accuracy'])

model_fash2.summary()
```

```
Model: "sequential_21"
```

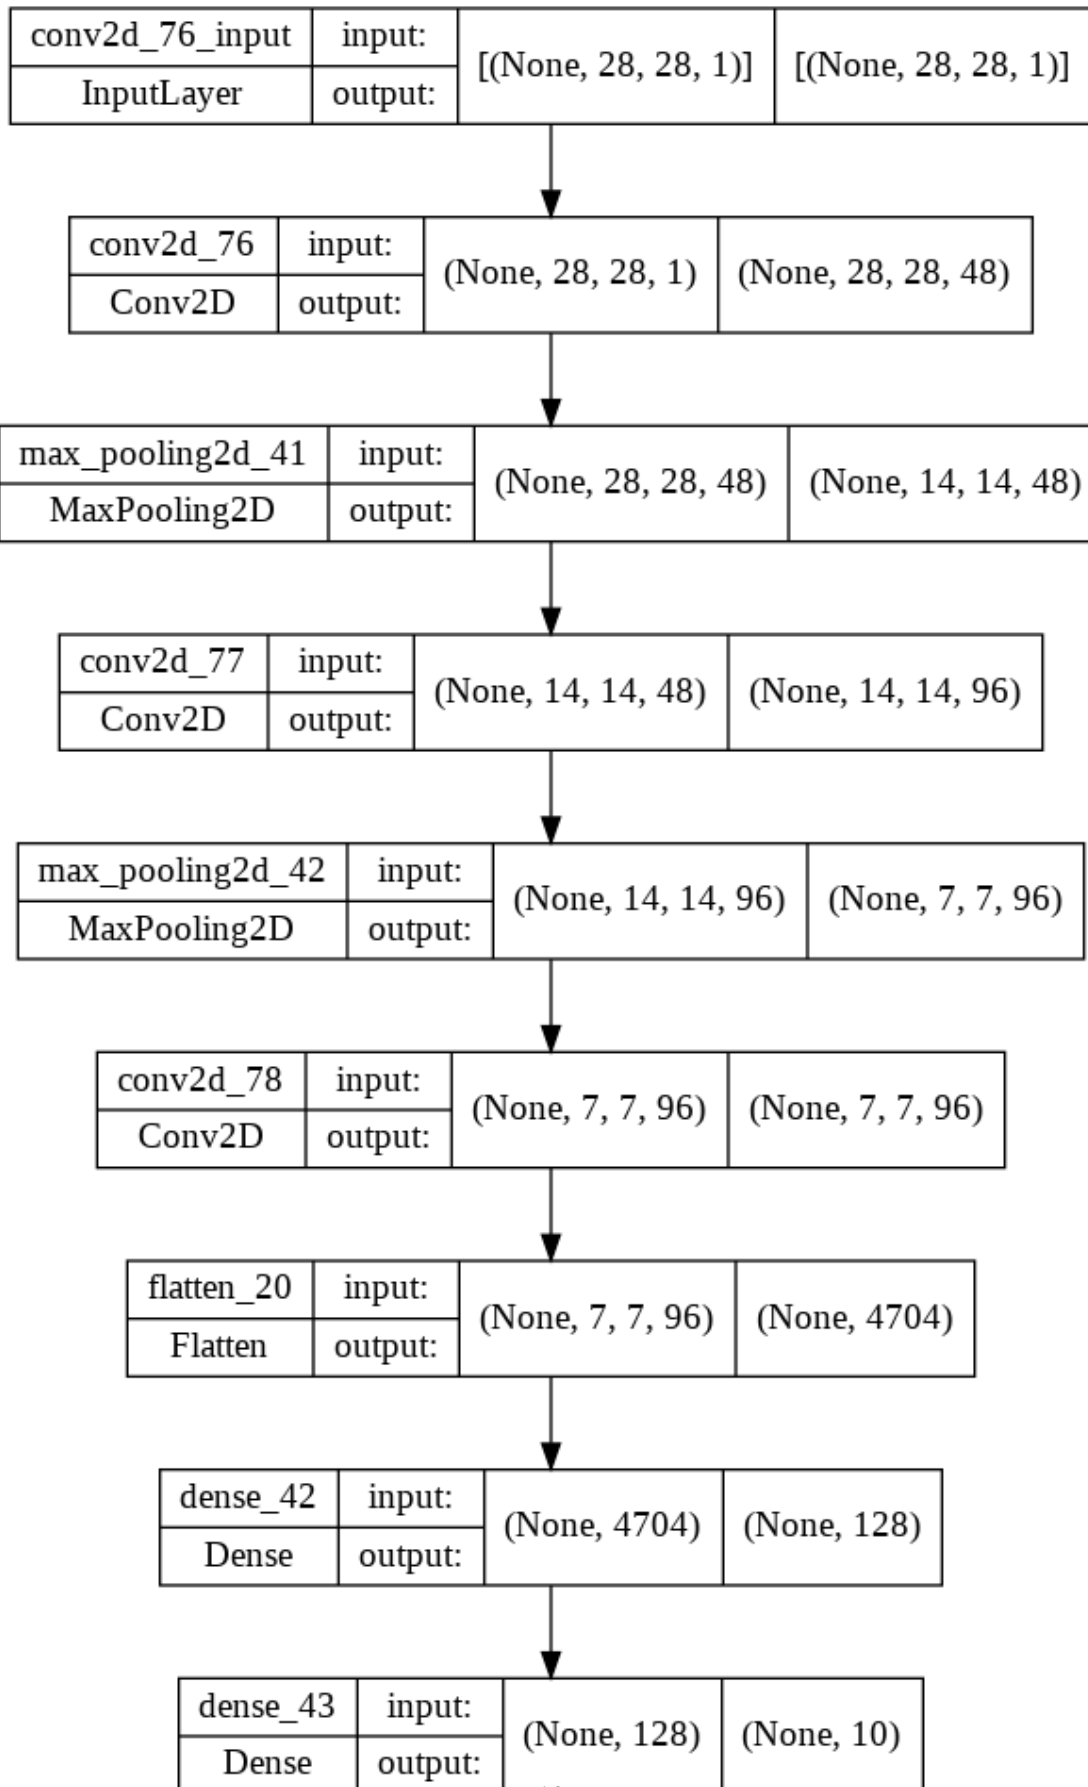
Layer (type)	Output Shape	Param #
conv2d_76 (Conv2D)	(None, 28, 28, 48)	480

max_pooling2d_41 (MaxPoolin g2D)	(None, 14, 14, 48)	0
conv2d_77 (Conv2D)	(None, 14, 14, 96)	41568
max_pooling2d_42 (MaxPoolin g2D)	(None, 7, 7, 96)	0
conv2d_78 (Conv2D)	(None, 7, 7, 96)	83040
flatten_20 (Flatten)	(None, 4704)	0
dense_42 (Dense)	(None, 128)	602240
dense_43 (Dense)	(None, 10)	1290

```
=====
Total params: 728,618
Trainable params: 728,618
Non-trainable params: 0
-----
```

```
[148]: tf.keras.utils.plot_model(model_fash2, 'model_fash2.png', show_shapes=True)
```

```
[148]:
```



```
[128]: model_fash2.fit(train_imgs_fash, train_labels_fash, epochs=15, batch_size=64)
```

Train on 60000 samples

Epoch 1/15

60000/60000 [=====] - 5s 86us/sample - loss: 0.4165 -
accuracy: 0.8475

Epoch 2/15

60000/60000 [=====] - 5s 81us/sample - loss: 0.2543 -
accuracy: 0.9079

Epoch 3/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.2126 -
accuracy: 0.9216

Epoch 4/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.1840 -
accuracy: 0.9335

Epoch 5/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.1622 -
accuracy: 0.9407

Epoch 6/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.1458 -
accuracy: 0.9474

Epoch 7/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.1315 -
accuracy: 0.9532

Epoch 8/15

60000/60000 [=====] - 5s 82us/sample - loss: 0.1223 -
accuracy: 0.9574

Epoch 9/15

60000/60000 [=====] - 5s 81us/sample - loss: 0.1139 -
accuracy: 0.9599

Epoch 10/15

60000/60000 [=====] - 5s 83us/sample - loss: 0.1042 -
accuracy: 0.9628

Epoch 11/15

60000/60000 [=====] - 5s 83us/sample - loss: 0.0975 -
accuracy: 0.9657

Epoch 12/15

60000/60000 [=====] - 5s 81us/sample - loss: 0.0946 -
accuracy: 0.9677

Epoch 13/15

60000/60000 [=====] - 5s 81us/sample - loss: 0.0911 -
accuracy: 0.9681

Epoch 14/15

60000/60000 [=====] - 5s 81us/sample - loss: 0.0862 -
accuracy: 0.9699

```
Epoch 15/15
60000/60000 [=====] - 5s 81us/sample - loss: 0.0837 -
accuracy: 0.9708
```

```
[128]: <keras.callbacks.History at 0x7f8517268310>
```

```
[129]: test_loss, test_acc = model_fash2.evaluate(test_imgs_fash, test_labels_fash)
print(f"Accuracy on the test dataset = {test_acc}")
```

```
/usr/local/lib/python3.7/dist-packages/keras/engine/training_v1.py:2057:
UserWarning: `Model.state_updates` will be removed in a future version. This
property should not be used in TensorFlow 2.0, as `updates` are applied
automatically.
```

```
updates = self.state_updates
```

```
Accuracy on the test dataset = 0.9185000061988831
```

6 Visualizing Filter Response

We use gradient descent in input space to display the visual pattern each filter is maximally responsive to.

To this end we take a VGG19 convnet pretrained on the ImageNet dataset.

Very Deep Convolutional Networks for Large-Scale Image Recognition Karen Simonyan, Andrew Zisserman

DL Architecture

```
[105]: from tensorflow.keras.applications import VGG19
from tensorflow.keras import backend as K
import numpy as np

import matplotlib.pyplot as plt

#Load pretrained model
#we omit the densely connected layers of the network
model = VGG19(weights='imagenet', include_top=False)
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-
applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80142336/80134624 [=====] - 0s 0us/step
80150528/80134624 [=====] - 0s 0us/step
```

```
[106]: model.summary()
```

```
Model: "vgg19"
```

```
-----
Layer (type)                Output Shape          Param #
=====
```

input_1 (InputLayer)	[(None, None, None, 3)]	0
block1_conv1 (Conv2D)	(None, None, None, 64)	1792
block1_conv2 (Conv2D)	(None, None, None, 64)	36928
block1_pool (MaxPooling2D)	(None, None, None, 64)	0
block2_conv1 (Conv2D)	(None, None, None, 128)	73856
block2_conv2 (Conv2D)	(None, None, None, 128)	147584
block2_pool (MaxPooling2D)	(None, None, None, 128)	0
block3_conv1 (Conv2D)	(None, None, None, 256)	295168
block3_conv2 (Conv2D)	(None, None, None, 256)	590080
block3_conv3 (Conv2D)	(None, None, None, 256)	590080
block3_conv4 (Conv2D)	(None, None, None, 256)	590080
block3_pool (MaxPooling2D)	(None, None, None, 256)	0
block4_conv1 (Conv2D)	(None, None, None, 512)	1180160
block4_conv2 (Conv2D)	(None, None, None, 512)	2359808
block4_conv3 (Conv2D)	(None, None, None, 512)	2359808
block4_conv4 (Conv2D)	(None, None, None, 512)	2359808
block4_pool (MaxPooling2D)	(None, None, None, 512)	0
block5_conv1 (Conv2D)	(None, None, None, 512)	2359808
block5_conv2 (Conv2D)	(None, None, None, 512)	2359808
block5_conv3 (Conv2D)	(None, None, None, 512)	2359808
block5_conv4 (Conv2D)	(None, None, None, 512)	2359808
block5_pool (MaxPooling2D)	(None, None, None, 512)	0

```

=====
Total params: 20,024,384
Trainable params: 20,024,384
Non-trainable params: 0

```

```

[107]: #Specify filter you want to visualize and get its output
layer_name = 'block5_conv3'
filter_index = 3
layer_output = model.get_layer(layer_name).output

#Loss is the averaged activation of the chosen filter
loss = K.mean(layer_output[:, :, :, filter_index])

[108]: #Gradients of loss with respect to the input
#upgrading to 2.x: tf.gradients is no longer supported
#requiring tf.compat.v1.disable_eager_execution()
grads = K.gradients(loss, model.input)[0]

#A trick is to normalize the gradients by their L2 norm
#This ensures that the magnitude of the gradients is always in the same range
#and leads to a smooth descent process
grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)

[109]: #The tensors defined so far (loss, grads) were symbolic
#To obtain values we need to feed an input via K.function

iterate = K.function([model.input], [loss, grads])
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])

[110]: # print(grads)
# print(grads_value)

[111]: #Implement the actual gradient descent
#Initial input is a grey image with some noise

input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step

[112]: # print(grads_value)

[116]: #Postprocess to turn into displayable image

def deprocess_image(x):
    x -= x.mean()
    x /= (x.std() + 1e-5)
    x *= 0.1

```

```


x += 0.5
x = np.clip(x, 0, 1)

x *= 255
x = np.clip(x, 0, 255).astype('uint8')
return x

```

```
[117]: plt.imshow(deprocess_image(input_img_data[0]))
```

```
[117]: <matplotlib.image.AxesImage at 0x7f85177af250>
```



HOMEWORK 2:

Write a function that takes as arguments the name of the layer and filter index and outputs the displayable filter response.

Then you can choose different filters and visualize which patterns they are responsive too! Submit the code (as Notebook) and at least 3 filter responses (a PDF file).

```

[143]: def visualize_layer(model, layer_name, filter_index):
        layer_output = model.get_layer(layer_name).output
        loss = K.mean(layer_output[:, :, :, filter_index])

        grads = K.gradients(loss, model.input)[0]
        #normalizing gradients by L2 norm
        grads /= (K.sqrt(K.mean(K.square(grads))) + 1e-5)

```

```


iterate = K.function([model.input], [loss, grads])
loss_value, grads_value = iterate([np.zeros((1, 150, 150, 3))])

input_img_data = np.random.random((1, 150, 150, 3)) * 20 + 128.
step = 1.
for i in range(40):
    loss_value, grads_value = iterate([input_img_data])
    input_img_data += grads_value * step


plt.imshow(deprocess_image(input_img_data[0]))
filename = layer_name + "_" + str(filter_index) + ".pdf"
plt.imsave(filename, deprocess_image(input_img_data[0]))

```

```
[144]: visualize_layer(model, 'block1_conv2', 63)
```



```
[145]: visualize_layer(model, 'block5_conv3', 500)
```

```
[146]: visualize_layer(model, 'block3_conv4', 200)
```

